

( / )

# Guide to Resilience4j

Last modified: May 6, 2020

by baeldung (<https://www.baeldung.com/author/baeldung/>)

**Cloud** (<https://www.baeldung.com/category/cloud/>)

**DevOps** (<https://www.baeldung.com/category/devops/>)

---

Get started with Spring 5 and Spring Boot 2, through the *Learn Spring* course:

**>> CHECK OUT THE COURSE** ([/ls-course-start](#))

---

## 1. Overview

In this tutorial, we'll talk about the Resilience4j (<https://resilience4j.github.io/resilience4j/>) library.

**The library helps with implementing resilient systems by managing fault tolerance for remote communications.**

The library is inspired by Hystrix (/introduction-to-hystrix) but offers a much more convenient API and a number of other features like Rate Limiter (block too frequent requests), Bulkhead (avoid too many concurrent requests) etc.

## 2. Maven Setup

To start, we need to add the target modules to our *pom.xml* (e.g. here we add the Circuit Breaker):

```
<dependency>
  <groupId>io.github.resilience4j</groupId>
  <artifactId>resilience4j-circuitbreaker</artifactId>
  <version>0.12.1</version>
</dependency>
```

Here, we're using the *circuitbreaker* module. All modules and their latest versions can be found on Maven Central (<https://search.maven.org/classic/#search%7Cga%7C1%7Cg%3A%22io.github.resilience4j%22>).

In the next sections, we'll go through the most commonly used modules of the library.


### 3. Circuit Breaker

Note that for this module we need the *resilience4j-circuitbreaker* dependency shown above.

The Circuit Breaker pattern (<https://martinfowler.com/bliki/CircuitBreaker.html>) helps us in preventing a cascade of failures when a remote service is down.

**After a number of failed attempts, we can consider that the service is unavailable/overloaded and eagerly reject all subsequent requests** to it. In this way, we can save system resources for calls which are likely to fail.

Let's see how we can achieve that with Resilience4j.



(<https://freestar.com/?>

First, we need to define the settings to use. The simplest way is to use default settings:

```
CircuitBreakerRegistry circuitBreakerRegistry
    = CircuitBreakerRegistry.ofDefaults();
```

It's also possible to use custom parameters:

```
CircuitBreakerConfig config = CircuitBreakerConfig.custom()
    .failureRateThreshold(20)
    .ringBufferSizeInClosedState(5)
    .build();
```

Here, we've set the rate threshold to 20% and a minimum number of 5 call attempts.

Then, we create a *CircuitBreaker* object and call the remote service through it:

```
interface RemoteService {
    int process(int i);
}

CircuitBreakerRegistry registry = CircuitBreakerRegistry.of(config);
CircuitBreaker circuitBreaker = registry.circuitBreaker("my");
Function<Integer, Integer> decorated = CircuitBreaker
    .decorateFunction(circuitBreaker, service::process);
```

Finally, let's see how this works through a JUnit test.



(<https://freestar.com/?>

We'll attempt to call the service 10 times. We should be able to verify that the call was attempted a minimum of 5 times, then stopped as soon as 20% of calls failed:

```
when(service.process(any(Integer.class))).thenThrow(new
    RuntimeException());

for (int i = 0; i < 10; i++) {
    try {
        decorated.apply(i);
    } catch (Exception ignore) {}
}

verify(service, times(5)).process(any(Integer.class));
```

### 3.1. Circuit Breaker's States and Settings

A *CircuitBreaker* can be in one of the three states:

- *CLOSED* – everything is fine, no short-circuiting involved
- *OPEN* – remote server is down, all requests to it are short-circuited

- *HALF\_OPEN* – a configured amount of time since entering OPEN state has elapsed and *CircuitBreaker* allows requests to check if the remote service is back online

We can configure the following settings:

- the failure rate threshold above which the *CircuitBreaker* opens and starts short-circuiting calls
- the wait duration which defines how long the *CircuitBreaker* should stay open before it switches to half open
- the size of the ring buffer when the *CircuitBreaker* is half open or closed
- a custom *CircuitBreakerEventListener* which handles *CircuitBreaker* events
- a custom *Predicate* which evaluates if an exception should count as a failure and thus increase the failure rate

## 4. Rate Limiter

Similar to the previous section, this features requires the *resilience4j-ratelimiter*

(<https://search.maven.org/classic/#search%7Cga%7C1%7Ca%3A%22resilience4j-ratelimiter%22>) dependency.

As the name implies, **this functionality allows limiting access to some service**. Its API is very similar to *CircuitBreaker's* – there are *Registry*, *Config* and *Limiter* classes.

Here's an example of how it looks:

```
RateLimiterConfig config =
    RateLimiterConfig.custom().limitForPeriod(2).build();
RateLimiterRegistry registry = RateLimiterRegistry.of(config);
RateLimiter rateLimiter = registry.rateLimiter("my");
Function<Integer, Integer> decorated
    = RateLimiter.decorateFunction(rateLimiter, service::process);
```

Now all calls on the decorated service block if necessary to conform to the rate limiter configuration.

We can configure parameters like:

- the period of the limit refresh
- the permissions limit for the refresh period
- the default wait for permission duration

## 5. Bulkhead

Here, we'll first need the *resilience4j-bulkhead* (<https://search.maven.org/classic/#search%7Cga%7C1%7Cresilience4j-bulkhead>) dependency.

It's possible **to limit the number of concurrent calls to a particular service**.

Let's see an example of using the Bulkhead API to configure a max number of one concurrent calls:

```
BulkheadConfig config =  
BulkheadConfig.custom().maxConcurrentCalls(1).build();  
BulkheadRegistry registry = BulkheadRegistry.of(config);  
Bulkhead bulkhead = registry.bulkhead("my");  
Function<Integer, Integer> decorated  
    = Bulkhead.decorateFunction(bulkhead, service::process);
```

To test this configuration, we'll call a mock service's method.

Then, we ensure that *Bulkhead* doesn't allow any other calls:

```
CountDownLatch latch = new CountDownLatch(1);
when(service.process(anyInt())).thenAnswer(invocation -> {
    latch.countDown();
    Thread.currentThread().join();
    return null;
});

ForkJoinTask<?> task = ForkJoinPool.commonPool().submit(() -> {
    try {
        decorated.apply(1);
    } finally {
        bulkhead.onComplete();
    }
});
latch.await();
assertThat(bulkhead.isCallPermitted()).isFalse();
```

We can configure the following settings:

- the max amount of parallel executions allowed by the bulkhead
- the max amount of time a thread will wait for when attempting to enter a saturated bulkhead

## 6. Retry

For this feature, we'll need to add the *resilience4j-retry* (<https://search.maven.org/classic/#search%7Cga%7C1%7Cresilience4j-retry>) library to the project.

We can **automatically retry a failed call** using the Retry API:

```
RetryConfig config = RetryConfig.custom().maxAttempts(2).build();
RetryRegistry registry = RetryRegistry.of(config);
Retry retry = registry.retry("my");
Function<Integer, Void> decorated
    = Retry.decorateFunction(retry, (Integer s) -> {
        service.process(s);
        return null;
    });
```

Now let's emulate a situation where an exception is thrown during a remote service call and ensure that the library automatically retries the failed call:

```
when(service.process(anyInt())).thenThrow(new RuntimeException());
try {
    decorated.apply(1);
    fail("Expected an exception to be thrown if all retries failed");
} catch (Exception e) {
    verify(service, times(2)).process(any(Integer.class));
}
```

We can also configure the following:

- the max attempts number
- the wait duration before retries
- a custom function to modify the waiting interval after a failure
- a custom *Predicate* which evaluates if an exception should result in retrying the call

## 7. Cache

The Cache module requires the *resilience4j-cache* (<https://search.maven.org/classic/#search%7Cga%7C1%7Cresilience4j-cache>) dependency.

The initialization looks slightly different than the other modules:

```
javax.cache.Cache cache = ...; // Use appropriate cache here
Cache<Integer, Integer> cacheContext = Cache.of(cache);
Function<Integer, Integer> decorated
    = Cache.decorateSupplier(cacheContext, () -> service.process(1));
```

Here the caching is done by the JSR-107 Cache (/jcache) implementation used and Resilience4j provides a way to apply it.



Note that there is no API for decorating functions (like *Cache.decorateFunction(Function)*), the API only supports *Supplier* and *Callable* types.

## 8. TimeLimiter

For this module, we have to add the *resilience4j-timelimiter* (<https://search.maven.org/classic/#search%7Cga%7C1%7Cresilience4j-timelimiter>) dependency.

It's possible to **limit the amount of time spent calling a remote service** using the *TimeLimiter*.

To demonstrate, let's set up a *TimeLimiter* with a configured timeout of 1 millisecond:

```
long ttl = 1;
TimeLimiterConfig config
    =
    TimeLimiterConfig.custom().timeoutDuration(Duration.ofMillis(ttl)).build
    ();
TimeLimiter timeLimiter = TimeLimiter.of(config);
```

Next, let's verify that Resilience4j calls *Future.get()* with the expected timeout:

```
Future futureMock = mock(Future.class);
Callable restrictedCall
    = TimeLimiter.decorateFutureSupplier(timeLimiter, () -> futureMock);
restrictedCall.call();

verify(futureMock).get(ttl, TimeUnit.MILLISECONDS);
```

We can also combine it with *CircuitBreaker*:

```
Callable chainedCallable
    = CircuitBreaker.decorateCallable(circuitBreaker, restrictedCall);
```

## 9. Add-on Modules

Resilience4j also offers a number of add-on modules which ease its integration with popular frameworks and libraries.

Some of the more well-known integrations are:

- Spring Boot – *resilience4j-spring-boot* module
- Ratpack – *resilience4j-ratpack* module
- Retrofit – *resilience4j-retrofit* module
- Vertx – *resilience4j-vertx* module
- Dropwizard – *resilience4j-metrics* module
- Prometheus – *resilience4j-prometheus* module


## 10. Conclusion

In this article, we went through different aspects of the Resilience4j library and learned how to use it for addressing various fault-tolerance concerns in inter-server communications.

As always, the source code for the samples above can be found over on GitHub (<https://github.com/eugenp/tutorials/tree/master/libraries-6>).

**Get started with Spring 5 and Spring Boot 2,  
through the *Learn Spring* course:**

>> **CHECK OUT THE COURSE** (</ls-course-end>)



## Build your Microservice Architecture with Spring Boot and Spring Cloud

**Download the E-book** (</spring-microservices-guide>)

Comments are closed on this article!

## COURSES

[ALL COURSES \(/ALL-COURSES\)](#)

[ALL BULK COURSES \(/ALL-BULK-COURSES\)](#)

[THE COURSES PLATFORM \(HTTPS://COURSES.BAELDUNG.COM\)](https://courses.baeldung.com)

## SERIES

[JAVA "BACK TO BASICS" TUTORIAL \(/JAVA-TUTORIAL\)](#)

[JACKSON JSON TUTORIAL \(/JACKSON\)](#)

[HTTPCLIENT 4 TUTORIAL \(/HTTPCLIENT-GUIDE\)](#)

[REST WITH SPRING TUTORIAL \(/REST-WITH-SPRING-SERIES\)](#)

[SPRING PERSISTENCE TUTORIAL \(/PERSISTENCE-WITH-SPRING-SERIES\)](#)

[SECURITY WITH SPRING \(/SECURITY-SPRING\)](#)

[SPRING REACTIVE TUTORIALS \(/SPRING-REACTIVE-GUIDE\)](#)

## ABOUT

[ABOUT BAELDUNG \(/ABOUT\)](#)

[THE FULL ARCHIVE \(/FULL\\_ARCHIVE\)](#)

[WRITE FOR BAELDUNG \(/CONTRIBUTION-GUIDELINES\)](#)

[EDITORS \(/EDITORS\)](#)

[JOBS \(/TAG/ACTIVE-JOB/\)](#)

[OUR PARTNERS \(/PARTNERS\)](#)

[ADVERTISE ON BAELDUNG \(/ADVERTISE\)](#)

[TERMS OF SERVICE \(/TERMS-OF-SERVICE\)](#)

[PRIVACY POLICY \(/PRIVACY-POLICY\)](#)

[COMPANY INFO \(/BAELDUNG-COMPANY-INFO\)](#)

[CONTACT \(/CONTACT\)](#)

